

Epic 2 Delivery Plan And Owner Questions

Date: 2026-07-03

Status: Planning task prepared; execution not started

Scope: E2 Research Engine To ICP

Executive Goal

Epic 2 converts the PRD Pack hypothesis into a source-graded market evidence engine before outreach, positioning upgrades, or demand claims.

The first goal is not to browse widely. The first goal is to define approved sources, data types, agent roles, scoring, competitor comparison, audience analysis, and owner decisions so real account research can proceed without contaminating the public repo or overstating evidence.

Current Starting Point

FACT:

- E2.0A is accepted.
- Primary ICP hypothesis is B2B SaaS product teams, Series B-D, 50-500 employees, 2-5 PMs, Director/VP Product owner.
- First offer hypothesis is the Product Discovery-to-Production PRD Pack.
- Approved source types are public company websites, public jobs, public reviews, public directories/lists, authorized LinkedIn summaries, approved search queries, and the internal approved corpus.
- Gated tools remain gated: Tavily, Exa, SerpAPI, Firecrawl, Apify, Playwright/browserless, Clay, Apollo, Cognism, Clearbit, People Data Labs.

GAP:

- No account universe has been built.
- No account evidence cards have real public signals.
- ICP trigger, budget path, workaround, and role currentness are still hypotheses.
- No 10-15 verified E6 outreach target set exists.

E2 Stage Plan

Stage	Output	Owner agents	Data/source types	Review gate
E2.1	Approved source list and segment data model	AF Tools, AF Research	Directories, company sites, job pages, internal corpus	Source boundary approval
E2.2	Competitor map and public pain evidence cards	AF Research	Competitor sites, pricing pages, reviews, public jobs, public case studies	A-D source grade
E2.3	One-page ICP card	AF Research, AF Review	Reviewed account cards and pain map	FACT/HYPOTHESIS/GAP split
E2.4	Primary ICP plus later-lane split	AF Manager, AF Review	Segment scoring and exclusions	No ICP drift

Stage	Output	Owner agents	Data/source types	Review gate
E2.5	10-15 verified targets for E6	AF Tools, AF Copy, AF Review	Company/domain/role-title records, no private personal data in public repo	Two B+ signals and role-currentness gate

Data Types To Collect

Data type	Source	Public repo handling
Company domain and name	Public directory/company site	Allowed
Segment and size estimate	Public directory/company site/careers page	Allowed with source grade
Product org signal	Careers page, public job post	Allowed, dated
PRD/discovery/handoff pain	Public review, blog, case study, job text	Summary only
Current tools/workarounds	Public tech/integration pages, reviews	Summary only
Buyer role title	Public leadership/team/careers pages or authorized LinkedIn summary	Role title only; no personal lead data by default
Social language	Public social posts	HYPOTHESIS only
Contact details	Gated enrichment/manual owner-approved workflow	Do not store in public repo

Competitor Analysis Plan

Initial comparison set:

- Productboard, Aha!, Dovetail, Zeda, BuildBetter, Fibery, Cycle, ChatPRD, Notion AI, Glean, Dust.

Compare on:

Field	Why it matters
Buyer and job	Shows whether they sell to the same owner
Artifact output	PRD, roadmap, insight repository, backlog, decision log
Evidence model	Whether claims are source-backed or AI-generated
Workflow integration	Jira, Linear, Notion, Productboard, Slack, docs
AI boundary	Where AI is used and what is reviewed
Pricing/CTA	Buying motion and budget proxy
Trust/privacy	Core objection for product-team source material
ArchFlow opportunity	The gap ArchFlow can own

Audience Analysis Plan

Analyze audience through four lenses:

Lens	Evidence	Output
Account shape	size, stage, product team size, market	segment fit
Trigger	hiring, funding, launch, research/rework language	timing score
Pain	PRD rework, discovery overload, handoff loss, roadmap ambiguity	pain type
Buying path	role owner, budget proxy, current workflow	outreach readiness

Tools And Sources

Use now:

- repo-local E2.0A schema and rubric;
- public company sites;
- public career pages and job boards;
- public review sites summarized only;
- public directories/lists;
- authorized LinkedIn summaries;
- manual approved search queries;
- dashboard/run packets after Codex review.

Use later only with approval:

- search APIs, crawlers, browser automation, enrichment vendors, provider-backed synthesis, live Notion/GitHub/TG writeback, and any personal-contact dataset.

Owner Questions Before Execution

- Should E2.1 use live public web/search research now, or stay local-plan-only?
- Which first source universe is approved: SaaS directories, funding lists, job boards, Product Hunt-style launch lists, or another list?
- Is authorized LinkedIn summary allowed for role-title currentness in this run?
- Should the first universe use 50-500 employees broadly or tighten to 100-300?
- What is the minimum evidence threshold for E6 handoff: score, two B+ signals, owner judgment, or all three?
- Should E3 diagnostic/analytics be finished before E4 social profile publication?
- Are any regions, industries, or account types excluded for the first pass?
- Is provider-backed summarization allowed after public-source collection, or must the first pass be deterministic/local only?

Planning Task Question Analysis

The E2 Notion task named PPlanning should contain both the questions and the analysis below. The default is conservative: do only public-source, public-safe, source-graded research until the owner explicitly approves live search, enrichment, provider-backed synthesis, or personal-contact handling.

Question	Why it matters	Default if unanswered	Risk if skipped
Live public web/search now or local-plan-only?	Decides whether E2 starts evidence collection or stays a method package	Local-plan-only plus approved source list	Accidental unbounded browsing or weak provenance
First source universe	Determines where accounts come from and what bias the sample has	SaaS directories plus public careers pages	Random account list with inconsistent fit
LinkedIn role-title currentness	Helps confirm buyer role without storing personal lead data	Authorized summaries only, role/title only	Stale or unverifiable buyer-owner assumptions
50-500 vs 100-300 employees	Controls segment breadth and product-team maturity	Keep 50-500, tag sub-segments	Overfitting too early or noisy ICP
Evidence threshold for E6	Prevents outreach based on weak signals	Require two independent B-or-better signals plus owner judgment	Low-quality outreach list

Question	Why it matters	Default if unanswered	Risk if skipped
E3 before E4 publication	Decides whether positioning and diagnostic should precede content	Finish E3 diagnostic/analytics before heavy publication	Content drives attention without conversion path
Exclusions	Prevents wasted research in bad-fit markets	Exclude regulated/high-security accounts until trust model matures	Privacy/security objections dominate too early
Provider summarization	Controls cost, privacy, and traceability	Deterministic/local first, provider only on sanitized digests after approval	Unlogged model use or contaminated source boundaries

Agent Distribution For Epic 2

Agent role	Responsibility	Required output
Source Boundary Agent	Define allowed public sources, forbidden data, and source-grade rubric	E2.1 source boundary packet
Universe Builder	Build the first company universe from approved public lists	Account universe table with source dates
Competitor Analyst	Compare alternatives and positioning gaps	Competitor comparison matrix
Audience Analyst	Score account shape, trigger, pain, and buying path	ICP signal cards
Evidence QA Agent	Grade provenance and separate fact/hypothesis/gap	Evidence-card audit notes
Offer Strategist	Translate evidence into PRD Pack offer, objections, and CTA	ICP card and offer hypothesis
Outreach Gatekeeper	Prepare E6 handoff only after evidence threshold passes	10-15 reviewed target candidates, no personal-contact dump

Competitor And Market Analysis Method

Use direct, dated public evidence first. Each competitor row should include the buyer, artifact workflow, AI/review boundary, integrations, pricing/CTA, proof assets, privacy/trust posture, and ArchFlow gap.

Compare competitors in three groups:

Group	Examples	What to learn
Product management suites	Productboard, Aha!, Zeda, Cycle	Existing budget category and integration expectations
Research/insight tools	Dovetail, BuildBetter, Notion AI, Glean	Evidence handling, synthesis claims, trust posture
AI artifact/workflow tools	ChatPRD, Dust, Fibery, internal AI assistants	Speed expectations and AI-review objections

Market And Audience Analysis Sources

Source type	Use	Handling rule
Company sites	Segment, product category, positioning, workflow hints	Summaries with dated URLs
Careers pages and public jobs	Product-team maturity, tooling, pain language	Job text summarized, no raw dumps
Review sites	Pain/workaround language	Short summaries only
Public pricing/CTA pages	Budget and buying motion proxy	Dated capture and interpretation
Public directories/lists	Universe construction	Keep list provenance and inclusion rule
Authorized LinkedIn summaries	Role-title currentness	Role/title only by default
Internal approved corpus	Offer and method alignment	Public-safe derived summaries

Epic 2 Execution Gate

E2 execution should start only after the Planning task records:

- approved source universe;
- allowed/disallowed data types;
- first segment boundaries;
- evidence threshold for E6 handoff;
- provider/use-of-search decision;
- excluded account classes;
- reviewer who can approve source cards.

Acceptance Criteria

- E2.1 source list and segment model are approved before collection.
- E2.2 evidence cards include provenance, grade, signal count, pain type, risk note, and next action.
- Shortlisted accounts have at least two independent B-or-better signals.
- E2.3 ICP card separates FACT, HYPOTHESIS, and GAP.
- E2.4 keeps later audiences separate.
- E2.5 hands E6 a reviewed 10-15 target set with no public overclaiming.
- Public safety scan passes before any repo-stored research packet is complete.

Notion Task Packet

Create or update the E2 Notion task page named Planning with this plan, then add the owner questions as the first execution checklist. Do not flip E2 tasks to Done. E2 should move from Backlog/To Do to Review/In Progress only after the owner answers the source and execution questions.