

Epic 2 Delivery Plan And Owner Questions

Date: 2026-07-03

Status: plan complete, execution not started

Scope: E2 Research Engine To ICP

Executive Goal

Epic 2 converts the PRD Pack hypothesis into a source-graded market evidence engine before outreach, positioning upgrades, or demand claims.

The first goal is not to browse widely. The first goal is to define approved sources, data types, agent roles, scoring, competitor comparison, audience analysis, and owner decisions so real account research can proceed without contaminating the public repo or overstating evidence.

Current Starting Point

FACT:

- E2.0A is accepted.
- Primary ICP hypothesis is B2B SaaS product teams, Series B-D, 50-500 employees, 2-5 PMs, Director/VP Product owner.
- First offer hypothesis is the Product Discovery-to-Production PRD Pack.
- Approved source types are public company websites, public jobs, public reviews, public directories/lists, authorized LinkedIn summaries, approved search queries, and the internal approved corpus.
- Gated tools remain gated: Tavily, Exa, SerpAPI, Firecrawl, Apify, Playwright/browserless, Clay, Apollo, Cognism, Clearbit, People Data Labs.

GAP:

- No account universe has been built.
- No account evidence cards have real public signals.
- ICP trigger, budget path, workaround, and role currentness are still hypotheses.
- No 10-15 verified E6 outreach target set exists.

E2 Stage Plan

Stage	Output	Owner agents	Data/source types	Review gate
E2.1	Approved source list and segment data model	AF Tools, AF Research	Directories, company sites, job pages, internal corpus	Source boundary approval
E2.2	Competitor map and public pain evidence cards	AF Research	Competitor sites, pricing pages, reviews, public jobs, public case studies	A-D source grade
E2.3	One-page ICP card	AF Research, AF Review	Reviewed account cards and pain map	FACT/HYPOTHESIS/GAP split
E2.4	Primary ICP plus later-lane split	AF Manager, AF Review	Segment scoring and exclusions	No ICP drift

Stage	Output	Owner agents	Data/source types	Review gate
E2.5	10-15 verified targets for E6	AF Tools, AF Copy, AF Review	Company/domain/role-title records, no private personal data in public repo	Two B+ signals and role-currentness gate

Data Types To Collect

Data type	Source	Public repo handling
Company domain and name	Public directory/company site	Allowed
Segment and size estimate	Public directory/company site/careers page	Allowed with source grade
Product org signal	Careers page, public job post	Allowed, dated
PRD/discovery/handoff pain	Public review, blog, case study, job text	Summary only
Current tools/workarounds	Public tech/integration pages, reviews	Summary only
Buyer role title	Public leadership/team/careers pages or authorized LinkedIn summary	Role title only; no personal lead data by default
Social language	Public social posts	HYPOTHESIS only
Contact details	Gated enrichment/manual owner-approved workflow	Do not store in public repo

Competitor Analysis Plan

Initial comparison set:

- Productboard, Aha!, Dovetail, Zeda, BuildBetter, Fibery, Cycle, ChatPRD, Notion AI, Glean, Dust.

Compare on:

Field	Why it matters
Buyer and job	Shows whether they sell to the same owner
Artifact output	PRD, roadmap, insight repository, backlog, decision log
Evidence model	Whether claims are source-backed or AI-generated
Workflow integration	Jira, Linear, Notion, Productboard, Slack, docs
AI boundary	Where AI is used and what is reviewed
Pricing/CTA	Buying motion and budget proxy
Trust/privacy	Core objection for product-team source material
ArchFlow opportunity	The gap ArchFlow can own

Audience Analysis Plan

Analyze audience through four lenses:

Lens	Evidence	Output
Account shape	size, stage, product team size, market	segment fit
Trigger	hiring, funding, launch, research/rework language	timing score
Pain	PRD rework, discovery overload, handoff loss, roadmap ambiguity	pain type
Buying path	role owner, budget proxy, current workflow	outreach readiness

Tools And Sources

Use now:

- repo-local E2.0A schema and rubric;
- public company sites;
- public career pages and job boards;
- public review sites summarized only;
- public directories/lists;
- authorized LinkedIn summaries;
- manual approved search queries;
- dashboard/run packets after Codex review.

Use later only with approval:

- search APIs, crawlers, browser automation, enrichment vendors, provider-backed synthesis, live Notion/GitHub/TG writeback, and any personal-contact dataset.

Owner Questions Before Execution

- Should E2.1 use live public web/search research now, or stay local-plan-only?
- Which first source universe is approved: SaaS directories, funding lists, job boards, Product Hunt-style launch lists, or another list?
- Is authorized LinkedIn summary allowed for role-title currentness in this run?
- Should the first universe use 50-500 employees broadly or tighten to 100-300?
- What is the minimum evidence threshold for E6 handoff: score, two B+ signals, owner judgment, or all three?
- Should E3 diagnostic/analytics be finished before E4 social profile publication?
- Are any regions, industries, or account types excluded for the first pass?
- Is provider-backed summarization allowed after public-source collection, or must the first pass be deterministic/local only?

Acceptance Criteria

- E2.1 source list and segment model are approved before collection.
- E2.2 evidence cards include provenance, grade, signal count, pain type, risk note, and next action.
- Shortlisted accounts have at least two independent B-or-better signals.
- E2.3 ICP card separates FACT, HYPOTHESIS, and GAP.
- E2.4 keeps later audiences separate.
- E2.5 hands E6 a reviewed 10-15 target set with no public overclaiming.
- Public safety scan passes before any repo-stored research packet is complete.

Notion Task Packet

Create or update the E2 Notion task page with this plan, then add the owner questions as the first execution checklist. Do not flip E2 tasks to Done. E2 should move from Backlog/To Do to Review/In Progress only after the owner answers the source and

execution questions.