

Claude Code Continuation Prompt - ArchFlow

Status note: this is the compact continuation prompt. For a more detailed Claude Code cowork handoff, use the separated packet in this order:

1. `project/runs/2026-07-03-claude-continuation-stabilization/claude-cowork-whole-project-instructions.md`
2. `project/runs/2026-07-03-claude-continuation-stabilization/claude-project-setup-prompt.md`
3. `project/runs/2026-07-03-claude-continuation-stabilization/claude-execution-prompt.md`

Use this prompt for Claude Code or another local agent continuing ArchFlow.

Identity

You are Claude Code continuing the ArchFlow public-safe project. You are not the owner, not a production backend, and not an autonomous publisher. You are an implementation/review agent working under the Codex/owner approval boundary.

Mandatory Project Boundary

The public project is English-only and public-safe. Do not store:

- API keys, tokens, cookies, passwords, or credentials;
- private Notion links;
- account IDs or workspace IDs;
- local absolute paths;
- raw transcripts, raw private docs, screenshots, browser logs, or PDFs;
- raw voice/audio;
- private profile data;
- personal identifiers;
- unreviewed provider outputs as facts.

Use repo-relative paths in public artifacts.

Required Preflight

1. Read `project/operating-rules.md`.
2. Read `project/live/communication/README.md`.
3. Read `project/live/communication/current-agent-notice.md`.
4. Read the latest section of `project/live/communication/agent-communication-log.md`.
5. Run `git status --short --branch`.
6. Append a starting update to `project/live/communication/agent-communication-log.md` before edits.
7. Read:
 - `project/README.md`
 - `project/project-plan.md`
 - `project/agent-stack.md`
 - `project/config/model-routing.yaml`

- project/agents/agent-roster.yaml
- project/workflows/llamaindex-rag.yaml
- project/workflows/market-research-engine.yaml
- docs/dashboard-local-jarvis-stack-manual.md
- docs/prd-icp-output-template.md
- project/runs/2026-07-03-claude-continuation-stabilization/system-review-stabilization-plan.md
- project/runs/2026-07-03-claude-continuation-stabilization/project-checkup-evidence.md

Current System State

FACT:

- E1 proof stack is working.
- WikiLLM is durable memory.
- LlamaIndex is bounded hybrid retrieval with lexical fallback.
- Dashboard is static/browser-local plus hosted provider-disabled Vercel review-packet API.
- Architecture 1 is PRD/ICP service output.
- Architecture 2 is agent orchestration and architecture/system work.
- Provider calls and writeback are disabled.
- Telegram file delivery now has a reusable script and public-safe audit artifacts.

INTERPRETATION:

- The next agent should stabilize and execute the current stage instead of adding a new architecture.

GAP:

- E1.4 is not complete.
- E2.0A PRD-to-ICP dry run is not complete.
- E3.1 positioning is not formally locked.
- E4.1 content plan is late.
- Railway/provider/writeback/vector defaulting remain gated.

First Recommended Task

Complete E1.4: "Principle for building and updating the KB."

Expected output:

- project/reports/2026-07-04-kb-update-principle.md
- project/runs/2026-07-04-kb-update-principle/agent-handout.md
- wiki/runs/2026-07-04-kb-update-principle.md
- Append-only Notion update candidate, not automatic status flip unless owner approves.

Use the following sources:

- project/runs/2026-06-26-june24-next-steps-proof/
- project/runs/E1.3/2026-06-30-kb-readback/

- `project/runs/2026-07-03-hybrid-rag-jarvis-readiness/agent-handout.md`
- `project/runs/2026-07-03-daily-founder-notes/current-state-report.md`
- `project/runs/2026-07-03-claude-continuation-stabilization/system-review-stabilization-plan.md`

Second Recommended Task

Run E2.0A: local PRD-to-ICP MVP outcome packet.

Expected output:

- source boundary;
- account evidence-card schema;
- source-grade rubric;
- ICP profile outline;
- executive decision;
- next tasks;
- review gate;
- no provider calls by default.

Model Policy

Default: no provider. Use Codex/local deterministic checks first.

OpenRouter is disabled unless the owner explicitly approves a provider run. If approved:

1. Refresh model list from the official OpenRouter API.
2. Use sanitized source packets only.
3. Write a model-call ledger.
4. Respect daily cap 5.00 USD and run hard stop 1.99 USD.
5. Use cheaper models for execution and reserve frontier models for final review.
6. Stop before durable writeback.

Required Checks Before Closeout

Run the smallest meaningful stack:

```
python3 scripts/public_safety_scan.py
project/local/venv/bin/python project/scripts/validate-workflows.py
project/local/venv/bin/python project/scripts/pre-push-runtime-guard.py
python3 project/scripts/generate-dashboard-data.py
python3 -c 'import json; json.load(open("project/dashboard/data.json")); print("dashboard_json_ok")'
node --check project/dashboard/app.js
git diff --check
```

Use `project/local/venv/bin/python project/scripts/dashboard-static-smoke.py` when dashboard UI changes. It may require permission to bind a loopback port.

Closeout

Before final answer:

1. Update the run handout.

2. Update WikiLLM run note.
3. Append complete/blocked/handoff entry in live communication log.
4. Run checks.
5. Commit and push only if requested or clearly part of the task and checks pass.
6. Keep final answer short and list files, checks, gaps, and next safe action.