

Claude Execution Prompt

Date: 2026-07-03

Status: copy-ready prompt for actual project work after setup readiness is complete.

Copy This Prompt Into Claude Code For Execution

You are Claude Code continuing ArchFlow after completing the setup/readiness pass. Work as a senior system architect/operator and pragmatic implementation agent. Your goal is to execute the next safe project task while preserving the knowledge system, public-safety boundary, dashboard truth, and live coordination discipline.

Do not restart from scratch. Use the existing project structure and run artifacts.

Execution Goal

Primary next task:

Complete E1.4: define the principle for building and updating the knowledge base.

Secondary task only after E1.4 is complete and validated:

Run E2.0A: create the local PRD-to-ICP dry-run outcome packet.

Do not start E3.1, E4.1, provider activation, Railway, vector defaulting, Nexus writeback, dashboard writeback, or autonomous external sync until the primary task is complete or explicitly reprioritized.

Required Context Files

Read these first:

```
project/operating-rules.md
project/live/communication/README.md
project/live/communication/current-agent-notice.md
project/live/communication/agent-communication-log.md
project/README.md
project/project-plan.md
project/agent-stack.md
project/config/model-routing.yaml
project/agents/agent-roster.yaml
project/workflows/knowledge-integration.yaml
project/workflows/llamaindex-rag.yaml
project/workflows/market-research-engine.yaml
wiki/index.md
wiki/memory.md
wiki/insights.md
wiki/log.md
project/runs/2026-07-03-claude-continuation-stabilization/claude-cowork-whole-project-instructions.md
project/runs/2026-07-03-claude-continuation-stabilization/claude-project-setup-prompt.md
project/runs/2026-07-03-claude-continuation-stabilization/system-review-stabilization-plan.md
project/runs/2026-07-03-claude-continuation-stabilization/project-checkup-evidence.md
project/runs/2026-07-03-hybrid-rag-jarvis-readiness/agent-handout.md
project/runs/2026-07-03-daily-founder-notes/current-state-report.md
```

Read these skills:

```
skills/arcagcom/SKILL.md
skills/task-handout/SKILL.md
skills/archflow-e1-runtime-guard/SKILL.md
skills/archflow1/SKILL.md
skills/archflow-task-breakdown/SKILL.md
skills/outquestions/SKILL.md
```

Before Edits

1. Run:

```
git status --short --branch
```

2. Check whether another agent has claimed the same files in the live communication log.

3. Append a starting entry to:

```
project/live/communication/agent-communication-log.md
```

4. Claim only the files you will edit. For E1.4, likely files are:

```
project/reports/<date>-kb-update-principle.md
project/runs/<date>-kb-update-principle/agent-handout.md
wiki/runs/<date>-kb-update-principle.md
wiki/log.md
project/dashboard/data.json
project/live/communication/agent-communication-log.md
```

5. If you need to change workflow or script files, announce scope expansion first.

E1.4 Expected Output

Create a clear public-safe KB update principle report, suggested path:

```
project/reports/<date>-kb-update-principle.md
```

The report must define:

- what counts as durable project knowledge;
- what stays as transient run evidence;
- what stays private and must never be ingested;
- how WikiLLM, LlamaIndex, Graphify, Notion, Obsidian/Nexus, dashboard packets, and Telegram evidence relate;
- how agents decide whether to write wiki/memory.md, wiki/insights.md, wiki/runs/, issues, or decisions;
- how to keep the KB useful for tracebacks without becoming noisy;
- how to keep source paths and provenance in retrieval outputs;
- how to handle wrong actions, bad captures, credential-boundary mistakes, or provider/runtime confusion;
- how to record FACT / INTERPRETATION / HYPOTHESIS / GAP;
- how to close a run with checks, handout, dashboard data, and live-log completion.

The report must be detailed enough that a new agent can update the KB correctly without reading chat history.

E1.4 Required Structure

Use this structure:

```
# KB Update Principle

Date:
Status:

## Executive Rule

## Layer Responsibilities

## What Becomes Durable Knowledge

## What Does Not Become Durable Knowledge

## Source And Provenance Rules

## Retrieval Rules

## WikiLLM Writing Rules

## Dashboard Packet Rules

## Notion And External Sync Rules
```

```
## Error And Traceback Rules

## Agent Workflow

## Validation

## Acceptance Criteria

## Next Actions
```

E1.4 Acceptance Criteria

The task is acceptable when:

- it preserves public-safe boundaries;
- it explains WikiLLM as durable memory and LlamaIndex as retrieval only;
- it keeps Graphify as generated structural reference;
- it keeps Nexus/Obsidian gated unless schema/tool proof exists;
- it explains when to write memory, insight, issue, decision, run, and log files;
- it defines how dashboard packets become durable knowledge only after agent review;
- it includes a traceback section for credential-boundary and wrong-capture corrections;
- it includes checks and next safe action;
- validation passes.

E2.0A Expected Output

Only start after E1.4 is complete.

Create a local PRD-to-ICP dry-run packet using approved public-safe sources only.

Expected artifacts:

```
project/runs/<date>-prd-icp-dry-run/source-boundary.md
project/runs/<date>-prd-icp-dry-run/account-evidence-card-schema.md
project/runs/<date>-prd-icp-dry-run/source-grade-rubric.md
project/runs/<date>-prd-icp-dry-run/icp-profile-outline.md
project/runs/<date>-prd-icp-dry-run/executive-decision.md
project/runs/<date>-prd-icp-dry-run/agent-handout.md
wiki/runs/<date>-prd-icp-dry-run.md
```

Use:

```
project/runs/2026-07-03-claude-continuation-stabilization/templates/prd-icp-service-template.md
project/workflows/market-research-engine.yaml
docs/prd-icp-output-template.md
```

No provider calls by default.

Model And Provider Rules

Default:

- do not call providers;
- use local deterministic scripts and source-grounded writing;
- use Codex/Claude as local operators only.

OpenRouter is allowed only if the owner explicitly approves a separate provider run. If approved:

1. Refresh model availability from the official model list.
2. Use sanitized source packets only.

3. Record a model-call ledger.
4. Keep browser JavaScript away from provider keys.
5. Enforce daily and per-run budget caps.
6. Assign reviewer.
7. Stop before durable writeback.

Do not use provider output as fact without source-grounded review.

MCP And Connector Rules During Execution

Before using a connector:

- list tools;
- inspect schemas;
- search/read before write;
- confirm target workspace/project/vault;
- record proof;
- avoid private URLs and IDs in tracked files.

During this execution:

- GitHub: local Git is enough unless issue/PR sync is explicitly requested.
- Notion: prepare update candidates; do not flip task state without owner approval.
- Obsidian/Nexus: do not use unless live schema/tool proof is explicitly requested.
- Telegram: do not send unless explicitly requested and using approved sender.
- OpenRouter: disabled unless explicitly approved.
- Railway/Vercel/Figma: not needed for E1.4.

Validation Commands

For E1.4 documentation work, run:

```
python3 scripts/public_safety_scan.py
python3 project/scripts/generate-dashboard-data.py
python3 -c 'import json; json.load(open("project/dashboard/data.json")); print("dashboard_json_ok")'
git diff --check
```

If you touch workflows, runtime, skills, or RAG, also run:

```
project/local/venv/bin/python project/scripts/validate-workflows.py
project/local/venv/bin/python project/scripts/pre-push-runtime-guard.py
project/local/venv/bin/python project/scripts/llamaindex-approved-corpus.py --mode smoke
project/local/venv/bin/python project/scripts/llamaindex-rag-benchmark.py
```

If you touch dashboard UI, also run:

```
node --check project/dashboard/app.js
project/local/venv/bin/python project/scripts/dashboard-static-smoke.py
```

Closeout Required

Before final response:

1. Update or create project/runs/<run-id>/agent-handout.md.
2. Update wiki/runs/<run-id>.md.
3. Append to wiki/log.md.

4. Append complete, blocked, or handoff to the live communication log.
5. Regenerate project/dashboard/data.json if public text changed.
6. Run validation.
7. Leave a concise final response with:
 - files changed;
 - checks run;
 - blockers/gaps;
 - next safe action.

Do not push unless explicitly approved for this execution.

Stop Conditions

Stop and ask for direction if:

- a task requires private raw source ingestion;
- an MCP write is needed but schema discovery is unavailable;
- a provider call is needed but no approval/ledger/budget exists;
- a public file would need private URLs, IDs, credentials, or local absolute paths;
- another agent has claimed the same files;
- validation fails and the fix would require broad refactor;
- the task drifts into Railway, provider-backed Jarvis, dashboard writeback, vector defaulting, or live Nexus without explicit approval.

Final Output Style

Use:

```
FACT:
- ...

INTERPRETATION:
- ...

GAP:
- ...

FILES:
- ...

CHECKS:
- ...

NEXT SAFE ACTION:
- ...
```

Keep chat concise. Put detailed durable content into files.