

ArchFlow System Review And Stabilization Plan

Date: 2026-07-03

Status: planning and handoff artifact for the next agent.

Executive Conclusion

ArchFlow is at a useful but fragile transition point. The E1 proof system works. The public-safe repository is structured. WikiLLM holds durable project memory. LlamaIndex is now a bounded hybrid retrieval layer with lexical fallback. The dashboard is a static/browser-local Jarvis control surface with hosted provider-disabled review-packet endpoints. The website now communicates the PRD/ICP service direction. Notion has a coherent E1-E7 board with Agent Tags, but several short-term tasks need sequence cleanup and owner acceptance.

The next agent should not add more broad architecture. The next agent should stabilize the current stage, then execute the first commercial-learning lane:

1. Finish E1.4 as the repeatable KB update principle.
2. Run E2.0A as a local PRD-to-ICP outcome packet from approved sources.
3. Tighten E3.1 positioning and make website claims match the accepted ICP.
4. Convert E4.1 into a five-week content calendar tied to measurable demand signals.
5. Close TG-1/TG-2/TG-3 with the new Telegram file-sender evidence and a redacted outbound template.

Company Goal

FACT:

- The current company direction is a Product Discovery-to-Production PRD Pack for product teams inside B2B SaaS scaleups.
- The operational method is to turn raw product-team material into source inventory, context digest, PRD, ICP logic, backlog, decision log, KB update, and review handoff.

INTERPRETATION:

- The company should be evaluated by whether real product leaders trust the system with source material, decision context, and budget authority.
- Internal proof is necessary but not sufficient. It only becomes business evidence when it produces buyer conversations, source-packet willingness, paid diagnostic interest, or firm paid intent.

GAP:

- No payment-test demand is validated yet.
- E2 account evidence cards and E6 outreach have not started at validated scale.

Current Stage

FACT:

- E1.1 is Done.
- E1.2 is in Review pending owner acceptance.
- E1.3 readback and RAG readiness are strong enough to support the next method-writing and research lanes.
- E1.4 is still Backlog and should be completed before widening research.
- E2.0A is the next safe strategic research dry run.

INTERPRETATION:

- The project should move from infrastructure proof to disciplined research/output proof.
- The dashboard and Jarvis are useful as control and reporting surfaces, not as autonomous execution yet.

Knowledge System Optimality

Current strengths

- WikiLLM is the correct durable memory layer.
- Live communication log prevents silent parallel-agent conflicts.
- Run handouts make future continuation practical.
- LlamaIndex now retrieves over approved corpus with stable chunk metadata and source paths.
- Graphify remains correctly scoped as generated structure, not final truth.
- Public-safety scan catches local paths, secrets, personal identifiers, and private links.

Current weaknesses

- E1.4 principle is missing, so the KB method is proven but not fully codified.
- Not all Notion tasks consistently have Agent Tags.
- There is still a risk that agents confuse dashboard packets with durable memory.
- Vector retrieval is not truly active because the local embedder is unavailable.
- Nexus/Obsidian live writeback is not proven in this public project lane.

Stabilization plan

1. Write E1.4 KB update principle as the canonical method.
2. Make every new run produce: run note, handout, source boundary, checks, and next safe action.
3. Keep LlamaIndex as retrieval only and WikiLLM as durable truth.
4. Add a vector-readiness run only after local embedder availability is proven.
5. Add Nexus/Obsidian only through explicit schema discovery and public-safe summaries.

Dashboard Optimality

FACT:

- Dashboard static smoke passed: 8 routes, provider calls 0, writeback 0.
- Hosted /health, /api/lanes/prd-icp, /api/lanes/agent-orchestra, and /api/chat return provider-disabled review packets.
- The dashboard has the correct two-lane split: Architecture 1 PRD/ICP and Architecture 2 Agent Orchestra.

INTERPRETATION:

- The dashboard is optimal for the current stage as an operator command center.
- It should not yet become the primary memory store, source editor, provider executor, or autonomous Notion writer.

GAP:

- Header API state is useful, but runtime state should be more explicit about local bridge, hosted provider-disabled API, Railway unavailable, and writeback disabled.
- Browser-local packets need a clearer export-to-Codex review flow.
- Dashboard-driven writeback should remain blocked until backend auth, source filters, and approval queue exist.

Next improvements:

1. Add a runtime state panel: static/local/hosted/Railway, provider mode, writeback mode, last health check, queue state.
2. Add packet export buttons that save review packets for Codex, not direct writes.
3. Add an E2.0A guided workflow: source-grade cards, ICP outline, evidence gaps, executive decision.
4. Add a TG status card after the Telegram file sender is accepted.

Website Optimality

FACT:

- Website is static and public-safe.
- Hero narrows to B2B SaaS product teams and PRD/ICP operating systems.
- It avoids claiming validated ROI or autonomous provider execution.
- Diagnostic page exists as a browser/local lead-capture concept.

INTERPRETATION:

- The website is good enough as a proof/positioning surface, but not yet a conversion system.
- It needs clearer proof artifacts, one primary CTA, and analytics instrumentation before traffic.

GAP:

- E3.1 positioning is still Late in Notion even though current source is close.
- Diagnostic form analytics and lead funnel are not proven.
- No payment-test evidence exists.

Next improvements:

1. Complete E3.1 as a formal positioning lock.
2. Add proof blocks that point to PRD/ICP outputs, review gates, and source-boundary discipline.
3. Connect E3.4 diagnostic form to E5 lead funnel metrics before traffic.
4. Keep SEO/noindex policy intentional until owner approves publication.

Notion Task Optimality

FACT:

- Task table has useful status, type, deadline, notes, GitHub link, and Agent Tags.
- Done items are generally evidence-backed.
- Backlog and Late items show the real next work rather than hiding gaps.

INTERPRETATION:

- The task board is useful, but the next PM pass should reduce duplicate or stale runtime tasks.
- TG-1/TG-2/TG-3 should be revisited after this run because Telegram file sending now has a script and audit artifacts.

GAP:

- No Notion task status was changed in this run.
- E1.2 owner acceptance remains unresolved.
- E1.4, E2.0A, E3.1, E4.1, and TG cleanup are the near-term bottlenecks.

Recommended short-term order:

1. E1.4: KB update principle.
2. E2.0A: PRD-to-ICP dry-run outcome packet.
3. E3.1: positioning lock.
4. E4.1: five-week content plan.
5. TG cleanup: sender contract, redacted template, health/audit proof.

Skill And Agent Setup For Claude Code

Claude Code should not invent a parallel operating system. It should adapt the project-local skills as operating contracts:

- `skills/task-handout/SKILL.md`: mandatory for continuation handoffs.
- `skills/arcagcom/SKILL.md`: live communication and file-claim discipline.
- `skills/archflow1/SKILL.md`: Jarvis/dashboard runtime gates.
- `project/agents/agent-roster.yaml`: role map and Agent Tags.
- `project/workflows/*.yaml`: execution contracts.
- `project/config/model-routing.yaml`: provider policy, not an active provider config.

If Claude has a local skill system, it should install or mirror only these project-local contracts first. It must not install network skills or external repos until a specific source boundary and safety gate exists.

Strategic Long-Term Plan

Phase A - Current stabilization

- Complete E1.4.
- Make the task board reflect the actual short-term sequence.
- Keep dashboard and website claims conservative.
- Keep provider and writeback disabled.

Phase B - First market evidence

- Run E2.0A and E2 evidence-card schema.
- Build 10 to 15 verified target accounts only inside the current ICP.
- Produce one ICP card and one outreach-ready offer packet.

Phase C - Conversion setup

- Lock E3.1 and E3.4.
- Instrument E5 funnel.
- Publish E4 content around proof, not broad AI claims.

Phase D - Outreach and verdict

- Run E6 wave with source-backed, low-friction messages.
- Record real conversations and source-packet willingness.
- Use E7 to judge payment or firm paid intent.

Deep Improvement Plan

1. Replace broad "AI operating system" wording with "Product Discovery-to-Production PRD Pack" in any current-facing surface that still drifts broad.
2. Create a canonical source-packet intake template.
3. Create a canonical PRD/ICP output packet template.
4. Add runtime-state dashboard panel.
5. Add E2.0A guided workflow to dashboard.
6. Add Telegram delivery audit page to dashboard data after the script is accepted.
7. Add OpenRouter activation checklist and ledger template before any provider call.
8. Keep vector/turbovec as a separate retrieval readiness lane.
9. Keep Nexus/Obsidian live writes separate from public repo RAG until schema discovery is proven.
10. Create a weekly PM closeout ritual: Notion task audit, WikiLLM run log, dashboard regen, checks, Git push, and Telegram summary.

Stop Rules For The Next Agent

Stop before:

- raw private ingestion;
- broad external repo ingestion;
- OpenRouter call without owner approval and ledger;
- Railway deployment;
- Notion Done-state promotion without evidence;
- dashboard-driven writeback;
- raw voice/audio storage;
- Telegram send without reviewed message and approved sender;
- public claim that demand is validated.